




NVIDIA®

AI FOR SCIENCE A PRACTICAL INTRODUCTION TO DEEP LEARNING

WITH KERAS AND TENSORFLOW



LEARNING GOALS




In couple of hours, we can only travel so far

Main goal:
Become familiar with main ideas and process

A starting point for solving your own problems




INTRO TO DL, PART 1




LAB 1: CNNs AND KERAS



LAB 2: TROPICAL CYCLONES



LAB 3: CFD STEADY FLOW




AGENDA


DEEP LEARNING ANALOGIES

What is this deep learning thing, anyway?

A NEW TYPE OF SOFTWARE




A GENERALIZATION OF CURVE FITTING



A NEW WAY TO BUILD SOFTWARE

Traditional Programming vs Machine Learning




A DIFFERENT WAY TO BUILD SOFTWARE

Traditional Programming vs Machine Learning


Goals for today:

1. Learn to use this new approach
2. Revolutionize Science



A DIFFERENT WAY TO BUILD SOFTWARE

Hand written vs learnt functions



HAND-WRITTEN FUNCTION

```
Function1(T,P,Q)  
update_mass()  
update_momentum()  
update_energy()  
do_mechanics()  
do_microphysics()  
y = get_precipitation()  
return y
```

LEARNED FUNCTION

```
Function1(T,P,Q)  
A = relu( w1 * [T,P,Q] + b1)  
B = relu( w2 * A      + b2)  
C = relu( w3 * B      + b3)  
D = relu( w4 * C      + b4)  
E = relu( w5 * D      + b5)  
y = sigmoid(w6 * E      + b6)  
return y
```

Convert expert
knowledge into a function

Reverse-engineer a function
from inputs / outputs



A DIFFERENT WAY TO BUILD SOFTWARE

The two approaches are complimentary



MANUAL PROGRAMMING

“SOFTWARE 1.0”
ENGINEERED
LABOR INTENSIVE
EXPLICIT
EXPLAINABLE
SIMPLE
FROM EXPERTISE

(For best results, combine as needed)


MACHINE LEARNING

“SOFTWARE 2.0”
REVERSE-ENGINEERED
AUTOMATIC
IMPLICIT
SUBTLE
COMPLEX
FROM EXAMPLES



A DIFFERENT WAY TO BUILD SOFTWARE

Complex phenomena are best described implicitly.




EXAMPLE: ATMOSPHERIC RIVER



A GENERALIZATION OF CURVE FITTING


Curve fitting provides the starting intuition




A GENERALIZATION OF CURVE FITTING

Differences from traditional curve fitting

Find f , given x and y



Supervised
Deep
Learning





IMPLEMENTATION BASICS


AUTO-ML

Eventually, the optimizer might be able to do everything for you



WHAT YOU NEED TO MAKE DEEP LEARNING WORK


You need three main ingredients (and some skill)



LARGE QUANTITIES OF DATA




ML FRAMEWORK



GPU ACCELERATOR


DEEP LEARNING FRAMEWORKS

Many frameworks to choose from (but not for Fortran)





DEVELOPMENT ENVIRONMENT

JUPYTER NOTEBOOKS



NVIDIA GPU CLOUD REGISTRY CONTAINERIZED SOFTWARE



LINEAR REGRESSION

With Scikit Learn


```
from sklearn.linear_model import LinearRegression
from numpy.random import rand
import numpy as np

# DATA
npts      = 50
x_train  = 10*rand(npts,1)
y_train  = 0.3*x_train + 0.5 * rand(npts,1)

# MODEL
model = LinearRegression()

# TRAIN
model.fit(x_train, y_train)

# TEST
x_test   = np.linspace(0,10,npts).reshape(-1,1)
y_predict = model.predict(x_test)
```



LINEAR REGRESSION

With TensorFlow and Keras


```
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import backend as K
from tensorflow.random import uniform
K.clear_session()

# DATA
npts      = 50
x_train   = 10*uniform(shape=[npts,1])
y_train   = 0.3*x_train + 0.5 * uniform(shape=[npts,1])

# MODEL
model = keras.models.Sequential()
model.add(keras.layers.Dense(1, input_shape=[1]))
optimizer = keras.optimizers.Adam(lr=1e-1)
model.compile(loss='mean_squared_error',optimizer=optimizer)

# TRAIN
hist = model.fit(x_train,y_train,epochs=100,verbose=0)

# TEST
x_test    = tf.linspace(0,10,npts)
x_test    = tf.reshape(x_test,[npts,1])
y_predict = model.predict(x_test)
```






TRAINING

TRAINING VS INFERENCE

TRAINING PHASE



SEARCH FOR THE RIGHT PIECES

ONLINE
LEARNING

INFERENCE PHASE




APPLY THE COMPLETED MODEL




TRAINING: THE PLAYERS

DATA, MODEL, LOSS, AND OPTIMIZER




TRAINING: GRADIENT DESCENT

Finding a solution is as easy as falling down a hill




OPTIMIZERS

Many variations on stochastic gradient descent



TRAINING: BACKPROPAGATION

Compute the gradient, by efficiently assigning blame



AUTOGRAD

Let a framework keep track of your gradient, so you don't have to


PyTorch Autograd

```
from torch.autograd import Variable


x = Variable(torch.randn(1, 10))
prev_h = Variable(torch.randn(1, 20))
W_h = Variable(torch.randn(20, 20))
W_x = Variable(torch.randn(20, 10))

i2h = torch.mm(W_x, x.t())
h2h = torch.mm(W_h, prev_h.t())
next_h = i2h + h2h
next_h = next_h.tanh()

next_h.backward(torch.ones(1, 20))
```




AI, MACHINE LEARNING, DEEP LEARNING



DEEP LEARNING VS. MACHINE LEARNING

When should I use deep learning vs traditional machine learning?



TRADITIONAL MACHINE LEARNING

Random forests, SVM, K-means, Logistic Regression
Features hand-crafted by experts
Small set of features: 10s or 100s
NVIDIA RAPIDS: orders of magnitude speedup

SUPERVISED DEEP LEARNING


CNN, RNN, LSTM, GAN, Variational Auto-encoders
Finds features automatically
High dimensional data: images, sounds, speech
Large set of labelled data (10k+ examples)
NVIDIA CU-DNN: accelerates DL frameworks




ARTIFICIAL NEURONS

Simple equations with adjustable parameters

Biological neuron



Artificial neuron



$$y = f(w_1x_1 + w_2x_2 + w_3x_3)$$

<https://towardsdatascience.com/the-differences-between-artificial-and-biological-neural-networks-a8b46db828b7>

CURVE FIT WITH SINGLE LAYER NEURAL NETWORK


```
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import backend as K
from tensorflow.random import uniform
K.clear_session()

# DATA
def f(x): return tf.math.sin(10*x)
x_train = uniform(shape=[100,1]) - 0.5
x_val   = uniform(shape=[ 20,1]) - 0.5
y_train = f(x_train)
y_val   = f(x_val)

# MODEL (using keras functional API)
inputs  = keras.Input(shape=[1])
x       = keras.layers.Dense(units=100,activation="relu")(inputs)
x       = keras.layers.Dense(1)(x)
outputs = keras.layers.Add()([inputs,x])
model   = keras.Model(inputs=inputs, outputs=outputs)
optimizer= keras.optimizers.Adam(lr=1e-2)
model.compile(loss='mean_squared_error',optimizer=optimizer)


# TRAIN
hist = model.fit(x_train,y_train,validation_data=(x_val,y_val),
                  epochs=5000,verbose=0, batch_size=1)

# TEST
x_test   = tf.linspace(-.5,.5,100)
y_predict = model.predict(x_test)
```




DATA SPLITTING

KEEP TEST, TRAINING, AND VALIDATION DATA SEPERATE




PLOTTING TRAINING AND VALIDATION LOSS

```
1 %matplotlib inline
2 import matplotlib.pyplot as plt
3
4 t_loss = hist.history['loss']
5 v_loss = hist.history['val_loss']
6 epoch = tf.range(1,len(t_loss)+1)
7 step = 200
|
9 plt.figure(figsize=(12,6),dpi=100)
10 plt.semilogy(epoch[::step],t_loss[::step],linewidth=2,label='training loss')
11 plt.semilogy(epoch[::step],v_loss[::step],linewidth=2,label='validation loss')
12 plt.xlabel('epoch'); plt.ylabel('loss'); plt.legend()
```



VISUALIZATION TOOLS: TENSORBOARD



```
tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir=".logs")
model.fit(x_train, y_train, epochs=2, callbacks=[tensorboard_callback])
# run the tensorboard command to view the visualizations.
```

```
tensorboard --logdir=path_to_your_logs
```






MODEL CAPACITY AND REGULARIZATION

MODEL CAPACITY

A good model is one that generalizes to new data



UNDERFIT

GOOD FIT

OVER FIT



GOOD-FIT

Checking for Generalization

Training Loop


```
xtrain, ytrain, xval, yval = load_data(npts=500, train_fraction = 0.10)

model      = taylor_series(order=5)
optimizer = torch.optim.Adam(model.params, lr=1.0e-2)
epochs     = 1000

for i in range(epochs+1):

    # training
    optimizer.zero_grad()
    yhat = model(xtrain)
    loss = (yhat - ytrain).pow(2).mean()
    loss.backward()
    optimizer.step()

    # validation
    yval_hat = model(xval)
    loss_val = (yval_hat - yval).pow(2).mean()
```



OVER-FITTING

Captures training data, but generalizes poorly

Training Loop

```
xtrain, ytrain, xval, yval = load_data(npts=500, train_fraction = 0.02)


model      = taylor_series(order=8)
optimizer  = torch.optim.Adam(model.params, lr=5.0e-3)
epochs     = 1000

for i in range(epochs+1):

    # training
    optimizer.zero_grad()
    yhat = model(xtrain)
    loss = (yhat - ytrain).pow(2).mean()
    loss.backward()
    optimizer.step()

    # validation
    yval_hat = model(xval)
    loss_val = (yval_hat - yval).pow(2).mean()
```

Use more data points
Reduce model capacity



UNDER-FITTING

Model is too simple to fit the curve

Training Loop

```
xtrain, ytrain, xval, yval = load_data(npts=500, train_fraction = 0.10)


model      = taylor_series order=2
optimizer = torch.optim.Adam(model.params, lr=1.0e-2)
epochs     = 1000

for i in range(epochs+1):

    # training
    optimizer.zero_grad()
    yhat = model(xtrain)
    loss = (yhat - ytrain).pow(2).mean()
    loss.backward()
    optimizer.step()


    # validation
    yval_hat = model(xval)
    loss_val = (yval_hat - yval).pow(2).mean()
```

Increase model capacity
Use a different model





REGULARIZATION

Early Stopping and Layer Regularization




```
callback = tf.keras.callbacks.EarlyStopping(monitor='loss', patience=3)  
callback = tf.keras.callbacks.EarlyStopping(monitor='loss', patience=3)
```




```
layer = tf.keras.layers.Dense(  
    5, input_dim=5,  
    kernel_initializer='ones',  
    kernel_regularizer=tf.keras.regularizers.L1(0.01),  
    activity_regularizer=tf.keras.regularizers.L2(0.01))
```

REGULARIZATION

BatchNorm and Dropout



```
model.add(Conv2D(60,3, padding = "same"))
model.add(BatchNormalization())
model.add(Activation("relu"))
```



```
model=keras.models.Sequential()
model.add(keras.layers.Dense(150, activation="relu"))
model.add(keras.layers.Dropout(0.5))
```

```
model.add(keras.layers.Dropout(0.2))
model.add(keras.layers.Dense(120, activation="relu"))
model.add(keras.layers.Softmax())
```







CHALLENGES AND POTENTIAL SOLUTIONS

LABELLING LARGE QUANTITIES OF DATA


How can we overcome the need for manual labelling?




Data Fusion
Using one data source
as the label for another



Self-Supervised Learning
Predicting input B from input A




Reinforcement Learning
Obtaining labels directly from the
environment or simulation




Human-in-the-loop
Using human machine iteration to
make labelling easier



TRANSFER LEARNING: DON'T START FROM SCRATCH



Train on simulated or related data




Fine-tune on the real data

ENFORCING PHYSICAL CONSTRAINTS




Conservation of Mass, Momentum, Energy, Incompressibility,
Turbulent Energy Spectra, Translational Invariance



Lagrange multipliers (penalization), Hard Constraints,
Projective Methods, Differentiable Programming

INTERPRETABILITY: EXPLAINABLE AI



Layer-wise Relevance Propagation






USING YOUR GPU

EXPLODING DATASETS

Logarithmic relationship between the dataset size and accuracy



Hestness, J., Narang, S., Ardalani, N., Diamos, G., Jun, H., Kianinejad, H., ... & Zhou, Y. (2017). Deep Learning Scaling is Predictable, Empirically. arXiv preprint arXiv:1712.00409.



THE SCALING LAWS

As you increase the dataset size you must increase the model size





Figure 1 Language modeling performance improves smoothly as we increase the model size, dataset size, and amount of compute² used for training. For optimal performance all three factors must be scaled up in tandem. Empirical performance has a power-law relationship with each individual factor when not bottlenecked by the other two.

Kaplan, J., McCandlish, S., Henighan, T., Brown, T. B., Chess, B., Child, R., ... & Amodei, D. (2020). Scaling Laws for Neural Language Models. *arXiv preprint arXiv:2001.08361*.




GPUS MAKE MACHINE LEARNING PRACTICAL

Train in a day, or a month?



ACCELERATED COMPUTING




PILLARS OF DATA SCIENCE PERFORMANCE

CUDA Architecture




Massively Parallel Processing

NVLink/NVSwitch



High Speed Connecting between
GPUs for Distributed Algorithms

CUDA-X AI



NVIDIA GPU Acceleration Libraries
for Data Science and AI




PILLARS OF DATA SCIENCE PERFORMANCE

CUDA Architecture




Massively Parallel Processing

NVLink/NVSwitch



High Speed Connecting between
GPUs for Distributed Algorithms



CUDA-X AI



NVIDIA GPU Acceleration Libraries
for Data Science and AI



TENSOR CORES FOR DIFFERENT NEEDS



- No Code Change Speed-up for Training


PILLARS OF DATA SCIENCE PERFORMANCE

CUDA Architecture




Massively Parallel Processing

NVLink/NVSwitch



High Speed Connecting between GPUs for Distributed Algorithms

CUDA-X AI



NVIDIA GPU Acceleration Libraries for Data Science and AI


PILLARS OF DATA SCIENCE PERFORMANCE

CUDA Architecture




Massively Parallel Processing

NVLink/NVSwitch



High Speed Connecting between
GPUs for Distributed Algorithms

CUDA-X AI




NVIDIA GPU Acceleration Libraries
for Data Science and AI



LEARNED FUNCTIONS ARE GPU ACCELERATED

Next level software. No porting required.




CONVOLUTION OPERATION

Pointwise multiply and sum, scalar output

$$int[c, p, q] = \sum_{i, j \in filter} Im[c][istart + i, jstart + j] . Filt[c][i, j]$$


$$output[p, q] = \sum_c int[c, p, q]$$




Input Image



Input Filter



Intermediate output



Final Output

Why do it once if you can do it n times ? Batch the whole thing.

VERIFYING YOUR GPU

TensorFlow

```
from tensorflow.python.client import device_lib  
print('GPU' in str(device_lib.list_local_devices()))
```

True

PyTorch

```
from torch import cuda  
print(f"cuda.is_available: {cuda.is_available()}")  
print(f"cuda.device_count(): {cuda.device_count()}")  
print(cuda.get_device_name(cuda.current_device()))
```

```
cuda.is_available: True  
cuda.device_count(): 2  
Quadro GV100
```

Keras

```
from keras import backend as K  
K.tensorflow_backend._get_available_gpus()
```

Using TensorFlow backend.

```
['/job:localhost/replica:0/task:0/device:GPU:0',  
 '/job:localhost/replica:0/task:0/device:GPU:1']
```

Julia

```
using CUDADrv, Printf  
@printf("device 0 = %s \n" , CUDADrv.name(CuDevice(0)))  
@printf("device 1 = %s \n" , CUDADrv.name(CuDevice(1)))
```

```
device 0 = Quadro GV100  
device 1 = Quadro GV100
```



TRAINING ON A SINGLE GPU

Keras

Automatically uses GPU if available

Julia

```
use_gpu = true
to_device(x) = use_gpu ? gpu(x) : x
x      = to_device.(x)
y      = to_device.(y)
model = to_device(model)
```

PyTorch

```
device = torch.device("cuda:1" if torch.cuda.is_available() else "cpu")

# move data onto the GPU
model = model.to(device)
X,Y   = X.to(device), Y.to(device)

# Alocate on the GPU directly
dtype = torch.cuda.FloatTensor
X = torch.zeros(100).type(dtype)

# set default location for tensors
torch.set_default_tensor_type('torch.cuda.FloatTensor')
```

NVIDIA-SMI

System Management Interface

✖ - □ Terminal File Edit View Search Terminal Help

Every 0.5s: nvidia-smi							Tue Sep 17 10:02:34 2019
Tue Sep 17 10:02:34 2019							
+-----+-----+-----+-----+-----+-----+						Driver Version: 396.82	
	NVIDIA-SMI 396.82						
	GPU Name Persistence-M Bus-Id Disp.A Volatile Uncorr. ECC						
	Fan Temp Perf Pwr:Usage/Cap Memory-Usage GPU-Util Compute M.						
=====+=====+=====+=====+=====+=====+=====+=====							
0 Quadro GV100 Off 00000000:04:00.0 On							Off
49% 61C P2 51W / 250W 1451MiB / 32500MiB							3% Default
-----+-----+-----+-----+-----+-----+-----+-----							
1 Quadro GV100 Off 00000000:84:00.0 Off							Off
41% 58C P2 116W / 250W 1701MiB / 32508MiB							41% Default
-----+-----+-----+-----+-----+-----+-----+-----							
+-----+-----+-----+-----+-----+-----+-----+-----+							
Processes: GPU Memory							
GPU PID Type Process name Usage							
-----+-----+-----+-----+-----+-----+-----+-----							
0 1661 G /usr/lib/xorg/Xorg 328MiB							
0 2381 G compiz 70MiB							
0 8598 G ...quest-channel-token=8542253777236559196 102MiB							
0 10245 C /home/dhall/anaconda3/envs/tf/bin/python 946MiB							
1 10245 C /home/dhall/anaconda3/envs/tf/bin/python 1690MiB							
-----+-----+-----+-----+-----+-----+-----+-----							

Memory Utilization

Processor Utilization

Process Info






FULLY CONNECTED NETWORKS (MULTI-LAYER PERCEPTRONS)

FULLY CONNECTED NETWORKS

A given neuron is connected to every neuron in the previous layer




ACTIVATION FUNCTIONS

Many to choose from. But most use ReLU or LeakyReLU


Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$




tanh

$$\tanh(x)$$




ReLU

$$\max(0, x)$$



Leaky ReLU

$$\max(0.1x, x)$$




Maxout

$$\max(w_1^T x + b_1, w_2^T x + b_2)$$

ELU

$$\begin{cases} x & x \geq 0 \\ \alpha(e^x - 1) & x < 0 \end{cases}$$



RELU BASIS FUNCTIONS

Piecewise continuous basis functions

Training Loop

```
import torch.nn as nn
xtrain, ytrain, xval, yval = load_data(npts=500, train_fraction = 0.30)

# MODEL
n_bases = 100
model = nn.Sequential(
    nn.Linear(1, n_bases),
    nn.ReLU(),
    nn.Linear(n_bases, 1))



optimizer = torch.optim.Adam(model.parameters(), lr=5.0e-3)
epochs = 6000

for i in range(epochs+1):

    # training
    optimizer.zero_grad()
    yhat = model(xtrain)
    loss = (yhat - ytrain).pow(2).mean()
    loss.backward()
    optimizer.step()

    # validation
    yval_hat = model(xval)
    loss_val = (yval_hat - yval).pow(2).mean()
```

(Pytorch Code)



MULTI-LAYER NETWORKS

Piecewise continuous basis functions

Training Loop

```
import torch.nn as nn

xtrain,ytrain,xval,yval=load_data(npts=500, train_fraction = 0.30)

# MODEL
n1,n2,n3 = 100,100,10

model = nn.Sequential(
    nn.Linear(1, n1), nn.ReLU(),
    nn.Linear(n1, n2), nn.ReLU(),
    nn.Linear(n2, n3), nn.ReLU(),
    nn.Linear(n3, 1))



optimizer = torch.optim.Adam(model.parameters(), lr=5.0e-3)

epochs      = 1000

for i in range(epochs+1):



    # training
    optimizer.zero_grad()
    yhat = model(xtrain)
    loss = (yhat - ytrain).pow(2).mean()
    loss.backward()
    optimizer.step()

    # validation
    yval_hat = model(xval)
    loss_val = (yval_hat - yval).pow(2).mean()
```



DEEPER NEURAL NETWORKS

More layers allows for more levels of abstraction



<https://web.eecs.umich.edu/~honglak/icml09-ConvolutionalDeepBeliefNetworks.pdf>



Large Scale Visual Recognition Challenge 2012



The Imagenet competition: Automatically classify images from 1000 different categories







CONVOLUTIONAL NEURAL NETWORKS

WHAT ARE CNNS USED FOR?


Problems with translational invariance



Computer Vision
Invariance in 2d space



Computational Physics
Invariance in 3d space



Audio and Time Series
Invariance in time



COMPUTER VISION TASKS

Each task requires a different model and data setup


Classification



Classification + Localization



Object Detection



Instance Segmentation









Image Credit: NERSC

CLASSIFICATION

Example: Classifying Land Use



Agricultural Airplane Baseballdiamond Beach Building Chaparral Denseresidential



Forest Freeway Golfcourse Harbor Intersection Mediumresidential Mobilehomepark



Overpass Parkinglot River Runway Sparseresidential Storage tanks Tenniscourt

UC Merced Land Use Database




ONE-HOT ENCODING

Input: Pixels, Output: One-hot encoding

INPUT:PIXEL VALUES



OUTPUT: ONE-HOT VECTOR




<https://blog.carboniteq.com/practical-image-recognition-with-tensorflow/>




IMAGES ARE POINTS, WITH MANY DIMENSIONS


IN: 3-D Vector



IN: 784-D Vector




FULLY CONNECTED NETWORKS AND IMAGES DON'T MIX




TRANSLATIONAL EQUIVARIANCE

Objects in nature look the same from place to place




WHAT IS A CONVOLUTION?

A small matrix transformation, applied at each point of the image



CONVOLUTION EXAMPLE: SOBEL FILTER




$$G_x = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}$$

$$G_y = \begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{bmatrix}$$

$$G = \sqrt{{G_x}^2 + {G_y}^2}$$



CONVOLUTION EXAMPLE: SOBEL FILTER



$$G_x = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}$$

$$G_y = \begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{bmatrix}$$


$$G = \sqrt{{G_x}^2 + {G_y}^2}$$






CLASSIFICATION

CLASSIFIER EVOLUTION OVER TIME




LENET-5

(1988) Yann LeCun. Hand written recognition. 60k parameters.



IMAGENET ILSVR COMPETITION

Large Scale Visual Recognition Competition (2010-2017)




<https://en.wikipedia.org/wiki/ImageNet>



ALEXNET

(2012): Krizhevsky, Sutskever, Hinton. ImageNet winner.

ImageNet Classification with Deep Convolutional Neural Networks




<https://en.wikipedia.org/wiki/AlexNet>




VGG-16

2014. ImageNet runner up. Simple, clean architecture.



INCEPTION-V1 (GOOGLENET)


2014. Train different size convolutions in parallel



A network graph visualization featuring numerous nodes represented by small circles. Some nodes are white, while others are a vibrant lime green. These nodes are interconnected by a dense web of thin, gray lines, forming a complex web-like structure. The overall effect is one of a large-scale, decentralized system or a social network.

RESNETS

MODELING TRENDS: DEEPER AND LARGER




Source: Source information is 14 pt, italic




PROBLEM: VANISHING GRADIENTS

Error signal decays exponentially as it propagates backward through the network

ERROR SIGNAL VANISHES DURING BACKPROP



DEEPER NETWORKS WERE HARDER TO TRAIN




<https://www.arxiv-vanity.com/papers/1512.03385/>




RESNETS AND SKIP CONNECTIONS

(aka Highway Networks)

ADD THE INPUT TO OUTPUT



DRAMATICALLY SIMPLIFIES THE LOSS LANDSCAPE



<https://arxiv.org/pdf/1512.03385.pdf>

<https://arxiv.org/abs/1712.09913>
https://jithinjk.github.io/blog/nn_loss_visualized.md.html




RESNET-50

2015 Microsoft Research. 50 Layers, 23M params.



Deep Residual Learning for Image Recognition

Kaiming He Xiangyu Zhang Shaoqing Ren Jian Sun
Microsoft Research




DENSENET


2017



INTRO TO DL, PART 1




INTRO TO DL, PART 2




LAB 1: CNNs AND KERAS



LAB 2: TROPICAL CYCLONES




LAB 3: CFD STEADY FLOW



LAB PART 1 CNN AND KERAS 101

MNIST

The standard ‘hello world’ problem for deep learning



MNIST

Keras implementation



```
1 from tensorflow import keras
2 from tensorflow.keras.datasets import mnist
3 from tensorflow.keras.models import Sequential
4 from tensorflow.keras.layers import Dense, Dropout, Flatten, Conv2D, MaxPooling2D
5 from tensorflow.keras import backend as K
6
7 num_classes = 10
8 img_rows, img_cols = 28, 28
9
10 # DATA
11 (x_train, y_train), (x_test, y_test) = mnist.load_data()
12
13 x_train = x_train.reshape(x_train.shape[0], img_rows, img_cols, 1)
14 x_test = x_test.reshape(x_test.shape[0], img_rows, img_cols, 1)
15 y_train = keras.utils.to_categorical(y_train, num_classes)
16 y_test = keras.utils.to_categorical(y_test, num_classes)
17
18 # MODEL
19 input_shape = (img_rows, img_cols, 1)
20 model = Sequential()
21 model.add(Conv2D(32, kernel_size=(3, 3), activation='relu', input_shape=input_shape))
22 model.add(Conv2D(64, (3, 3), activation='relu'))
23 model.add(MaxPooling2D(pool_size=(2, 2)))
24 model.add(Dropout(0.25))
25 model.add(Flatten())
26 model.add(Dense(128, activation='relu'))
27 model.add(Dropout(0.5))
28 model.add(Dense(num_classes, activation='softmax'))
29 model.compile(loss=keras.losses.categorical_crossentropy,
30                 optimizer=keras.optimizers.Adadelta(),
31                 metrics=['accuracy'])
32
33 # TRAIN
34 model.fit(x_train, y_train, batch_size=128, epochs=12,
35             verbose=1, validation_data=(x_test, y_test))
36
37 # TEST
38 score = model.evaluate(x_test, y_test, verbose=0)
39 print('Test loss:', score[0])
39 print('Test accuracy:', score[1])
```


FASHION MNIST

A slightly more interesting version of MNIST

Label	Description	Examples
0	T-Shirt/Top	
1	Trouser	
2	Pullover	
3	Dress	
4	Coat	
5	Sandals	
6	Shirt	
7	Sneaker	
8	Bag	
9	Ankle boots	



6 STEPS APPROACH



LAUNCH CNN PRIMER AND KERAS 101

12:30-1:00 ET


Step through the primer on your own (shift + enter on each cell)

The following contents will be covered during the Bootcamp :

- [CNN Primer and Keras 101 \(Intro to DL/Part 2.ipynb\)](#)
- [Tropical Cyclone Intensity Estimation using Deep Convolution Neural Networks.](#)

CLICK HERE

Shutdown the kernel before clicking on “Next Notebook” to free up the GPU memory



FORGOT TO SHUTDOWN YOUR KERNELS?

Don't worry, you can fix it.

`NotFoundError: Failed to get convolution algorithm. This is probably because cuDNN failed to initialize, so try looking to see if a warning log message was printed above.`

Go to Home Tab, Click Running Tab, Kill notebooks you aren't using


The screenshot shows the Jupyter Home interface. At the top, there are several tabs: 'My Int' (closed), 'Home' (highlighted with a red box), 'Start_...', 'Part_2...', and 'CNN's...'. Below the tabs is a toolbar with icons for back, forward, search, and other functions. The main area is titled 'jupyter' and contains three tabs: 'Files', 'Running' (highlighted with a red box), and 'Clusters'. Under 'Running', it says 'Currently running Jupyter processes' and shows a list of terminals: 'There are no terminals running.' In the 'Notebooks' section, three notebooks are listed: '/Start_Here.ipynb' (Python 3), '/Intro_to_DL/Part_2.ipynb' (Python 3), and '/Intro_to_DL/CNN's.ipynb' (Python 3). Each notebook entry has a 'Shutdown' button.

Restart & Clear Output on the Kernel you are using

The screenshot shows the Jupyter kernel menu. The 'Kernel' tab is selected. A dropdown menu is open, listing several options: 'Interrupt', 'Restart', 'Restart & Clear Output' (highlighted with a red box), 'Restart & Run All', 'Reconnect', 'Shutdown', and 'Change kernel'. Below the menu, the title 'CNN Primer' is visible, along with the text 'This notebook covers introduction to Convolutional Ne terminologies.'




LEVELS OF AI ENGAGEMENT



COMPUTATIONAL SCIENCES


Similarities to the shift
Feature → Network Engineering?

NNs as a Porting Strategy?



CAN THIS WORK ∀? ABSOLUTELY, YES!

Proof: Universal Approximation Theorem



Problem: this is an essentially useless theorem for practical purposes



WILL THIS WORK \forall ?


Considering pesky practical constraints, like memory and performance

- Anecdotal Evidence: \exists scientific cases where NNs seem to do work extremely well
- Save bet: it will not work for \forall
- Therefore, by induction (sort of):
 - There exists \exists a subspace in \forall HPC applications, for which AI works well
 - Need to explore the **size** and **shape** of this subspace
 - Currently I think it is fair to say we don't understand this domain very well
 - **But:** Each individual case promising 10x, 100x, 1000x performance improvement is probably worth exploring; those can be groundbreaking!

But Intuition is Misleading




WHAT MAKES AI * HPC SPECIAL?




HOW TO FILL IN THE



Guided Design



New Approaches



SCIENTIFIC CHALLENGES

Barriers to acceptance of deep learning as a tool for science

- **Interpretability:** Can I understand what the neural-net is doing?
- **Robustness:** Will it always give me the right answer?
- **Coverage:** How much training data do I need?
- **Convergence:** How can I ensure that training will converge?
- **Uncertainty:** How certain can I be of the answers?




ESTIMATING TROPICAL CYCLONE INTENSITY

Paper Overview


Tropical Cyclone Intensity Estimation Using a Deep Convolutional Neural Network

Ritesh Pradhan, Ramazan Aygun, *Senior Member, IEEE*, Manil Maskey, *Member, IEEE*, Rahul Ramachandran, *Senior Member, IEEE*, and Daniel Cecil



OUTPUT: 8 CLASSES


Category	Symbol	Wind speeds	Damage
Five	H5	\geq 137 knots	Catastrophic
Four	H4	113- 136 knots	Catastrophic
Three	H3	96- 112 knots	Devastating
Two	H2	83- 95 knots	Extensive
One	H1	64- 82 knots	Significant
Tropical storm	TS	34- 63 knots	Significant
Tropical depression	TD	20- 33 knots	Small
No Category	NC	\leq 20 knots	-




ESTIMATING TROPICAL CYCLONE INTENSITY

Background: Dvorak technique

Dvorak Technique (1974)



Advanced Dvorak Technique- version 9 (2019)




[https://doi.org/10.1175/1520-0493\(1975\)103%3C0420:TCIAAF%3E2.0.CO;2](https://doi.org/10.1175/1520-0493(1975)103%3C0420:TCIAAF%3E2.0.CO;2)

<https://doi.org/10.1175/WAF-D-19-0007.1>




ESTIMATING TROPICAL CYCLONE INTENSITY

CNN Model




6 STEPS APPROACH

Steps to follow while solving a Machine Learning problem



DATA



SAFFIR-SIMPSON HURRICANE WIND SCALE
AND RELATED CLASSIFICATIONS

Category	Symbol	Wind speeds	Damage
Five	H5	≥ 137 knots	Catastrophic
Four	H4	113–136 knots	Catastrophic
Three	H3	96–112 knots	Devastating
Two	H2	83–95 knots	Extensive
One	H1	64–82 knots	Significant
Tropical storm	TS	34–63 knots	Significant
Tropical depression	TD	20–33 knots	Small
No Category	NC	≤ 20 knots	-



TASK

Multi-class Classification.

NC (No Category , ≤ 20 knots)

TD (Tropical Depression , 20-33 knots)

TS (Topical Storm , 34-63 knots)

H1 (Category One , 64-82 knots)

H2 (Category Two , 83-95 knots)

H3 (Category Three , 96-112 knots)

H4 (Category Four , 113-136 knots)

H5 (Category Five , ≥ 137 knots)



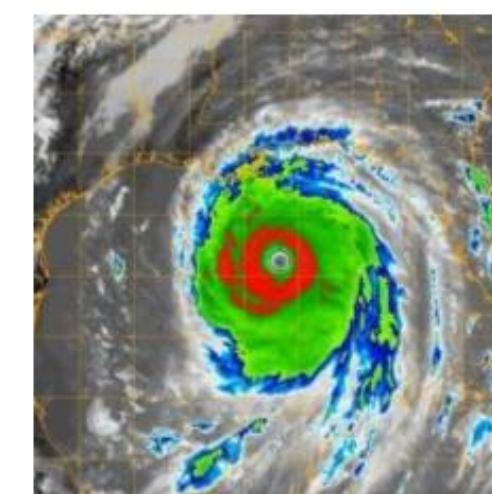
H2



H4



H3



H5



MODEL 1




Fig. 2. Network architecture for hurricane intensity estimation showing different steps of convolution and pooling.


Loss Function: Multi-class Cross-Entropy loss functions

Optimizer SGD (Stochastic Gradient Descent)

Training and Evaluation: Training Set 72 % , Test Set, 8 %, Validation Set 10%



SUMMARY OF APPROACH



PREPROCESSING DATA

Pre-Processing Data:

Step 1 : Resize Image from (1024, 1024 ,3) to (256 , 256 ,3)

Step 2 : Choose a random (232 , 232 , 3) patch from the (256 , 256 , 3) and feed into our model.

There are different types of Resizing:

- cv2.INTER_AREA (Preferable for Shrinking)
- cv2.INTER_CUBIC (Preferable for Zooming but slow)
- cv2.INTER_LINEAR (Preferable for Zooming and the default option)



LAUNCH TROPICAL CYCLONE NOTEBOOK

1:00-2:00 ET


Step through the notebook on your own (shift + enter on each cell)

The following contents will be covered during the Bootcamp :

- CNN Primer and Keras 101 (Intro to DL/Part 2.ipynb)
- Tropical Cyclone Intensity Estimation using Deep Convolution Neural Networks.

CLICK HERE

Shutdown the kernel before clicking on “Next Notebook” to free up the GPU memory



TROPICAL CYCLONE COMPETITION

Can you make a better prediction?

Go to last notebook in the TC section

The following contents will be covered during the Bootcamp :

- CNN Primer and Keras 101 (Intro to DL/Part 2.ipynb)
- Tropical Cyclone Intensity Estimation using Deep Convolution Neural Networks.

See if you can improve the accuracy.

Suggestions:

- Improve the data balance
- Tweak the model hyperparameters
- Try different optimizers


Bug in the lab: when doing the test/train split on time-series data, the data should not be shuffled!

- Try maximizing validation accuracy on both shuffled and un-shuffled validation sets



STEADY STATE FLOW WITH NEURAL NETWORKS

Flow fields are simulated using computational fluid dynamics (CFD) solvers




STEADY STATE FLOW WITH NEURAL NETWORKS

Our aim is to predict 2D flow around objects. The input is the boundary around which we want to calculate the flow. Here is an example of input data and the corresponding flow that was calculated using the Lattice Boltzmann method. ([Mechsys](#)).




6 STEPS APPROACH

Steps to follow while solving a Machine Learning problem



predict the velocity vectors of both the xx and yy channels from our model.

DATA AND TASK



MODEL

We will be building the following Models and benchmarking them as we proceed :

- Simple Fully Connected Networks

3 Layer Network

5 Layer Network

- Convolution Neural Networks

Binary Boundary

Signed Distance Function


- Advanced Networks

Gated Residual Network

Non-Gated Residual Network





SUMMARY OF APPROACH




U GATED NETWORK

2D Flow Prediction Network






AI FOR SCIENCE BOOTCAMP USING SIMNET




OUTLINE

- Introduction to Physics Informed Neural Networks (PINN)
- Solving Partial Differential Equation system using SimNet toolkit
- Solving parameterized PDEs
- Solving transient problems
- Solving inverse problems
- Challenge CFD problem - Flow over a 2D chip



DATA DRIVEN METHODS

Pros
Not dependent on Physics
Cons
No physics awareness; Generalization ability may be limited
Need to generate a lot of simulations (accuracy dependent on the simulation code)
Not very efficient for complex 3D geometries/curved surfaces
Interpolation/extrapolation errors



NEURAL NETWORK SOLVER THEORY

Goal: Train a neural network to satisfy the boundary conditions and differential equations by constructing an appropriate loss function

- Consider an example problem:
- We construct a neural network $u_{net}(x)$ which has a single value input $x \in \mathbb{R}$ and single value output $u_{net}(x) \in \mathbb{R}$.
- We assume the neural network is infinitely differentiable $u_{net} \in C^\infty$ - Use activation functions that are infinitely differentiable




NEURAL NETWORK SOLVER THEORY

- Construct the loss function. We can compute the second order derivatives $\left(\frac{\delta^2 u_{net}}{\delta x^2}(x) \right)$ using Automatic differentiation
- Where x_i are a batch of points in the interior $x_i \in (0, 1)$. Total loss becomes $L = L_{BC} + L_{Residual}$
- Minimize the loss using optimizers like Adam



NEURAL NETWORK SOLVER THEORY

- For $f(x) = 1$, the true solution is $\frac{1}{2}(x - 1)x$. After sufficient training we have,




SOLVING PARAMETERIZED PROBLEMS

- Consider the parameterized version of the same problem as before. Suppose we want to determine how the solution changes as we move the position on the boundary condition $u(l) = 0$
- Parameterize the position by variable $l \in [1, 2]$ and the problem now becomes:
- This time, we construct a neural network $u_{net}(x, l)$ which has x and l as input and single value output $u_{net}(x, l) \in \mathbb{R}$.
- The losses become



SOLVING PARAMETERIZED PROBLEMS

- For $f(x) = 1$, for different values of l we have different solutions





SOLVING INVERSE PROBLEMS

- For inverse problems, we start with a set of observations and then calculate the causal factors that produced them
- For example, suppose we are given the solution $u_{true}(x)$ at 100 random points between 0 and 1 and we want to determine the $f(x)$ that is causing it
- Train two networks $u_{net}(x)$ and $f_{net}(x)$ to approximate $u(x)$ and $f(x)$




SOLVING INVERSE PROBLEMS

- For $u_{true}(x) = \frac{1}{48} \left(8x(-1 + x^2) - \frac{3 \sin(4\pi x)}{\pi^2} \right)$ the solution for $f(x)$ is $x + \sin(4\pi x)$




SOLUTION TO PDES- 1D DIFFUSION







- Composite bar with material of conductivity $D_1 = 10$ for $x \in (0,1)$ and $D_2 = 0.1$ for $x \in (1,2)$. Point A and C are maintained at temperatures of 0 and 100 respectively
- Equations: Diffusion equation in 1D
- Flux and field continuity at interface ($x = 1$)




SOLUTION TO PDES- 1D DIFFUSION






- Define the problem and train the neural network to obtain the temperature distribution in the bar
- Compare the results with analytical solution




SOLUTION TO PARAMETERIZED PDES- 1D DIFFUSION



- Composite bar with material of conductivity D_1 for $x \in (0,1)$ and $D_2 = 0.1$ for $x \in (1,2)$.
- Solve the problem for multiple values of D_1 in the range (5, 25) in a single training
- Same boundary and interface conditions as before




SOLUTION TO ODES- COUPLED SPRING MASS SYSTEM







- Three masses connected by four springs
- System's equations (ordinary differential equations):
- For given values masses, spring constants and boundary conditions




SOLUTION TO ODES- COUPLED SPRING MASS SYSTEM





- Define the transient problem for time, $t = (0, 10)$ and train the neural network to obtain the displacement of each mass
- Compare the results with analytical solution




INVERSE PROBLEMS- COUPLED SPRING MASS SYSTEM



- For the same system, assume we know the analytical solution which is given by:
- With the above data and the values for m_2, m_3, k_1, k_2, k_3 same as before, use the neural network to find the values of m_1 and k_4



CHALLENGE: FLOW OVER 2D CHIP








- Solve the flow over 2D chip for the given boundary conditions. The challenge problem has 3 parts:
 1. Solve the fluid flow for the given boundary conditions and geometry
 2. Solve the fluid flow for the parameterized Chip geometry
 3. Solve the inverse problem where, given a flow field, use it to invert out the viscosity of the flow








CHALLENGE: FLOW OVER 2D CHIP

h: 0.6, w: 1.0



h: 0.4, w: 1.4



CHALLENGE: FLOW OVER 2D CHIP- HINTS AND TIPS

- Use Signed Distance Function to weight the equation losses inside domain for faster convergence (User Guide Section 2.3.2)
- Use Integral Continuity for faster convergence (User Guide Section 8.3.1 and 8.3.2)




SIMNET FEATURES AND ADVANCEMENTS

Physics types:





- Linear Elasticity (plane stress, plane strain and 3D)
- Fluid Mechanics
- Heat Transfer
- Coupled Fluid-Thermal
- Electromagnetics
- 2D wave propagation

Solution of differential equations:

- Ordinary Differential Equations
- Partial Differential Equations
- Differential (strong) Form
- Integral (weak) form of the PDEs



Taylor-Green vortex decay



Wave Propagation




SIMNET FEATURES AND ADVANCEMENTS

Several neural network architectures:


- Fully connected Network
- Fourier Feature Network
- Sinusoidal Representation Network (SiReN)
- Modified Fourier Network
- Deep Galerkin Method Network
- Modified Highway Network
- Multiplicative Filter Networks

Other Features include:


- Global and local learning rate annealing
- Global adaptive activation functions
- Halton sequences for low-discrepancy point cloud creation
- Gradient Accumulation
- Time-stepping schemes for transient problems
- Temporal loss weighting and time marching for the continuous time approach
- Importance sampling
- Homoscedastic task uncertainty quantification for loss weighting



Comparisons of various networks in SimNet applied to solve the flow over a heatsink





Importance sampling



Handling larger batch sizes using Multi-GPU and/or Gradient Aggregation

SIMNET FEATURES AND ADVANCEMENTS

- APIs to automatically generate point clouds from Boolean compositions of geometry primitives or import point cloud for complex geometry (e.g., STL files)
- Parameterized system representation that solves several configurations concurrently for analytical geometry using SimNet CSG module
- Transfer learning for efficient surrogate-based parameterization of STL and constructive solid geometries
- Polynomial Chaos Expansion method for assessing how uncertainties in a model input manifest in its output




SIMNET FEATURES AND ADVANCEMENTS


Improved performance with XLA enabled for TensorFlow models and multi-GPU/multi-Node runs

- Accelerated Linear Algebra (XLA)
- Strong scaling with learning rate adjustments

Improved stability in multi-GPU/multi-Node implementations using linear-exponential learning rate and utilization of TF32 precision for A100 GPUs




(a) Time per iteration




(b) Speed-up

Accelerated training using TF32 on A100 GPUs



Strong Scaling- Speedup and Scaling Efficiency





Thanks!

